

Relaxed Radix Balanced Trees

Elias Khoury

2016

Background

- Invented by researchers at EPFL in Lausanne
- RRB-Trees are part of the ongoing development of the Scala Programming Language
- The goal is to present a datastructure which can be concatenated in $O(\log n)$ time at no cost to other operations

Application

- RRB-Trees can be used to implement an efficient vector datastructure
- As it is a purely functional datastructure the operations on it can be easily run in parallel
- An application for efficient vector concatenation is:

Main Concept

- To represent a vector as a tree structure
- Multiple variables to control:
 - ▶ Branching factor
 - ▶ Balance of the tree
- The main operations are:
 - ▶ Indexing
 - ▶ Updating
 - ▶ Append to the front or back
 - ▶ Splitting

The Problem

Given two trees of branching factor m , it is impossible to merge the trees faster than linear time without degenerating to a linked list

Radix Search

- With a perfectly balanced tree of branching factor m , for any particular sub branch, there are exactly $m^{\hat{h}-1}$ leaves. Therefore, an index i can be found using $\lfloor i / (m^{\hat{h}-1}) \rfloor$

Relaxed Radix Search

- However, for this particular application we need to relax the branching factor m in order to circumvent the linear concatenation cost.
- The concept of a variable branching factor has been applied to tree structures before, like 2-3 Trees and finger trees
- By relaxing the branching factor, we lose the guarantee that any leaf node can be found in $O(1)$ time, this is due to the possibility of the tree being slightly unbalanced.
- The cost this adds is a slight linear lookup when arriving at a node that is unbalanced.

Branching Factor and Balance

- BALANCE

RRB-Trees

- The concatenation operation works by effectively merging the right side of the left tree, with the left side of the right tree.
- starting from the leaf nodes of both trees, a new tree is built from the bottom up containing the merged data of both trees.
- The concept behind this is to use the concept of sharing when building this new tree in order to minimise the cost of copying data
- For a tree of height, h . The complexity of this operation is $O(\log(h))$. However, there is a significant constant time cost when it comes to ensuring that the merged tree is balanced

Concatenation Process

A note on branching factor

- For simplicity of the diagrams a branching factor of 4 was chosen
- In practice the most efficient branching factor is actually 32
- Branching factor controls tree height which affects the runtime of the operations
- Need to minimise branching factor without degenerating into a linked list

	RRB-Vector	With $m = 32$
indexing	\log_m	eC
update	aC	eC
insert ends	$m \times \log_m$	aC
concat	$m^2 \times \log_m$	L v.s. eC
split	$m \times \log_m$	eC

Runtime Complexity In Relation To Similar Structures